



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

FTI: Fault Tolerance Interface

Presented by: Kai Keller

Developed by Dr. Leonardo Bautista Gomez



Motivation

Motivations

- Fault tolerance is critical at extreme scale.
- More components, more failures.
- Dense architectures, correlated failures.
- Multiple different types of failures (hard, soft, ect).
- Power limits might impact reliability.
- Current techniques will not scale.



Basic informations about FTI

- Download at <http://www.github.com/leobago/fti>
- Documentation at <http://leobago.github.io/fti>
- Library in c/c++ with Fortran bindings
- More than 8000 lines of code
- Applications ported:
 - HACC
 - Nek5K
 - CESM (ice module)
 - LAMMPS
 - GYSELA 5D
 - SPECFEM3D (CUDA version)
 - HYDRO
 - Other miniApps



Why FTI?

Multilevel Checkpointing

Local Storage: SSD, PCM, NVM.

Fastest checkpoint level.

Low reliability, transient failures.

Partner Copy: Ckpt. Replication.

Fast copy to neighbor node.

It tolerates single node crashes.

RS Encoding: Ckpt. Encoding.

Slow for large checkpoints.

Very reliable, multiple node crashes.

File System: Classic Ckpt.

Slowest of all levels.


The most reliable. Power outage.

- **Multiple:**


- Resiliency levels
- Checkpoint overheads
- Checkpoint intervals
- Power consumptions

Topology aware clustering

- FTI dedicated threads
- Asynchronous data transfer
- Reduced ckpt. Overhead
- Fine-grained control
- Disjoint communicators



Topology aware clustering



- Automatic process location recognition
- Intelligent clustering
- Enhanced reliability for node crashes
- Automatic repositioning after failure

A photograph of a large server rack filled with numerous hard drives. The rack is made of metal and has a glass front panel. The drives are arranged in a grid pattern, with some drives having labels on them. The background is dark, making the silver metal of the rack stand out.

How to use?

Easy-to-use API

- **Functions:**
 - `FTI_Init()`
 - `FTI_Protect()`
 - `FTI_Snapshot()`
 - `FTI_Finalize()`
- **Communicator:**
 - `FTI_COMM_WORLD`

```
int main(int argc, char **argv) {  
  
    MPI_Init(&argc, &argv);  
    FTI_Init("conf.fti", MPI_COMM_WORLD);  
  
    double *grid;  
    int i, steps=500, size=10000;  
    initialize(grid);  
    FTI_Protect(0, &i, 1, FTI_INTG);  
    FTI_Protect(1, grid, size,FTI_DFLT);  
  
    for (i=0; i<steps; i++) {  
        FTI_Snapshot();  
        kernel1(grid);  
        kernel2(grid);  
        comms(FTI_COMM_WORLD);  
    }  
  
    FTI_Finalize();  
    MPI_Finalize();  
    return 0;  
}
```

FTI_Init()

- **FTI_Init(confFile, communicator):**
 - Read/parse configuration file
 - Recognizes whether is a restart or not
 - Creates checkpoint directories
 - Detect topology of the system
 - Regenerates/moves data upon recovery
 - Splits the communicator (optional)

FTI_Protect()

- **FTI_Protect(ID,pointer,size,type):**
 - Stores metadata of the protected variable
 - FTI can predict size of checkpoints
 - Useful for data compression/aggregation
 - Can be reseted during the execution
 - User can create new FTI types
 - Required in order to write/read ckpt. data

FTI_Snapshot()

- **FTI_Snapshot():**
 - Measures (global average) iteration length
 - Exponential decay for global agreement
 - Translates from minutes to iterations
 - Test if it is time for a checkpoint
 - If it is, it checks which level of ckpt.
 - It saves the checkpoint as requested
 - It loads the checkpoint upon recovery
 - Planning to integrate notifications

Beyond FTI_Snapshot

- **FTI_Checkpoint(*ID*, *lvl*):**
 - Takes a checkpoint with id *ID* and level *lvl*
- **FTI_Status()**
 - Returns the status (initial run or restart)
- **FTI_Recover():**
 - It recovers from last available checkpoint

FTI_Finalize()

- **FTI_Finalize():**
 - Frees the allocated memory
 - Informs it is over to dedicated threads
 - Clean checkpoints and metadata
 - Moves last ckpt to PFS if requested

Configuration file for (1/3)

```
[basic]
# Set to 1 for having 1 FTI dedicated process per node
Head = 1

# Number of processes per node (including FTI dedicated processes)
node_size = 2

# Path where local checkpoints will be stored
ckpt_dir = /path/to/local/storage/

# Path where global checkpoints will be stored
lbl_dir = /path/to/global/storage/

# Path where checkpoints metadata will be stored
meta_dir = /path/to/myhome/.fti/

# Checkpoint interval in minutes for level 1
ckpt_int = 1

# Checkpoint interval in minutes for level 2
ckpt_12 = 2

# Checkpoint interval in minutes for level 3
ckpt_13 = 4

# Checkpoint interval in minutes for level 4
ckpt_14 = 8
```

Configuration file for FTI (2/3)

```
[basic]

# Set to 0 to do L2 post-processing asynchronously by the
# dedicated process
inline_12 = 0

# Set to 0 to do L3 post-processing asynchronously by the
# dedicated process
inline_13 = 0

# Set to 0 to do L4 post-processing asynchronously by the
# dedicated process
inline_14 = 0

# Set to 1 to keep the last checkpoint after Finalize
keep_last_ckpt = 0

# Size of the group for RS-encoding and Partner-copy ring
group_size = 4

# Set to 1 for verbose mode, 2 for moderate, 3 for silent
verbosity
```

Configuration file for (3/3)

```
[restart]

# This will be set to 1 automatically after FTI_Init
Failure          = 0

# This will be set to 1 automatically after FTI_Init
exec_id          = 2013-11-20_15-01-52

[advanced]

# Block size for communications
block_size        = 1024

# MPI tag for FTI communications
mpi_tag           = 2612

# Set to 1 for local tests in one single node
local_test        = 0
```

A photograph of a large server rack filled with multiple hard drive units. The rack is made of metal and has a perforated front panel. The word "Performance" is overlaid in white text.


Performance

Scaling to ~10K processes

- CURIE supercomputer in France
- SSD on the compute nodes (16 cores)
- HYDRO scientific application
- Using 1 FTI dedicated process per node
- Checkpointing every ~6 minutes
- Weak scaling to almost 10k processes

Scaling to ~10K processes

Weak Scaling Checkpointing Overhead



Scaling to >32K processes

- MIRA supercomputer at ANL (BG\Q)
- Persistent memory compute nodes
- LAMMPS scientific application
- Lennard-Jones simulation of 1.3 billion atoms
- 512 nodes, 64 MPI processes per node (32,678pr.)
- Using 1 FTI dedicated process per node
- Power monitoring during the entire run
- Checkpointing every ~5 minutes
- Less than 5% overhead on time to completion


Scaling to >32K processes

Synchronous Checkpointing

Without FTI - dedicated process

Head = 0

Execution: ~ 3000s



Scaling to >32K processes


Asynchronous Checkpointing

With FTI-dedicated process

Head = 1

Execution: ~ 2400s

10 minutes faster!



Features / Limitations

Interesting features

- FTI can predict time and size of next checkpoints
- Detailed knowledge of the datasets allows for transparent data compression/verification
- Transparent dedicated processes (Comm. Split)
- Topology reconstruction upon restart
- Dynamic checkpoint interval adaptation

Limitations

- FTI needs every rank in the given communicator to write a checkpoint file
- Application level checkpoint (code modification)
- Coordinated checkpoint, everybody restarts

A photograph of a large server rack filled with blue server units. The rack is made of metal and has glass doors. The servers are arranged in a grid pattern. The background is dark.

Thank you!